AIGP Module 1: Foundations of Artificial Intelligence

Comprehensive Study Notes

---

1. INTRODUCTION TO ARTIFICIAL INTELLIGENCE

1.1 What is Artificial Intelligence?

Definition: AI is the theory and development of computer systems able to perform tasks that normally require human intelligence.

Key Characteristics:

- Performs tasks requiring human-like intelligence
- Includes visual perception, speech recognition, decision-making, language translation
- Ranges from narrow/specific tasks to broader cognitive functions
- Continuously evolving field with expanding capabilities

Historical Context:

- Coined as a term in the 1950s
- Has experienced multiple waves of development ("AI winters" and "AI summers")
- Current era driven by increased computing power, big data, and advanced algorithms

---

2. THE TURING TEST

2.1 Concept and Purpose

Definition: A test proposed by Alan Turing in 1950 to evaluate a machine's ability to exhibit intelligent behavior indistinguishable from a human.

How It Works:

- A human evaluator engages in natural language conversations
- One conversation is with a human, another with a machine
- If the evaluator cannot reliably distinguish the machine from the human, the machine is said to have passed the test

Significance:

- Early benchmark for measuring AI intelligence
- Focused on conversational ability and human-like responses
- Sparked philosophical debates about machine consciousness and intelligence
- Still referenced today, though limitations are recognized

Narahari

Limitations:

- Focuses only on conversational ability, not general intelligence
- Doesn't test other forms of intelligence (visual, spatial, emotional)
- Can potentially be "fooled" by clever programming without true understanding

---

3. AI AS A SOCIO-TECHNICAL SYSTEM

3.1 Understanding Socio-Technical Systems

Definition: AI systems are socio-technical systems because they involve both **technical components** (hardware, software, algorithms) and **social components** (people, organizations, processes, societal impacts).

Why This Matters for Governance:

- AI doesn't operate in isolation—it affects and is affected by human behavior
- Technical solutions alone are insufficient; governance must address social impacts
- Requires interdisciplinary approaches (technical, legal, ethical, social)
- Changes in either technical or social elements affect the entire system

Components of AI as Socio-Technical System:

Technical Components:

- Algorithms and models
- Data infrastructure
- Computing hardware
- Software platforms
- Network architecture

Social Components:

- Developers and data scientists
- End users and affected populations
- Organizations deploying AI
- Regulatory bodies
- Societal norms and expectations

Interactions:

- Technical design choices reflect social values
- User behavior affects system performance

Narahari

- Organizational policies shape AI deployment
- Societal context influences acceptance and regulation

---

4. CHARACTERISTICS OF AI REQUIRING GOVERNANCE

4.1 Complexity

Definition: AI systems, especially modern ones, are often highly complex with intricate architectures and millions/billions of parameters.

Why It Requires Governance:

- Difficult to understand how decisions are made
- Hard to predict behavior in all situations
- Challenges in debugging and troubleshooting
- Requires specialized expertise to evaluate
- Creates accountability gaps when errors occur

Governance Implications:

- Need for technical documentation and transparency
- Importance of testing across diverse scenarios
- Requirements for expert oversight and review
- Documentation of system architecture and decision logic

4.2 Opacity (Black Box Problem)

Definition: Many AI systems, particularly deep learning models, operate as “black boxes” where the internal decision-making process is not transparent or easily interpretable.

Why It Requires Governance:

- Users and affected individuals can’t understand why a decision was made
- Difficult to identify bias or errors in reasoning
- Challenges in providing meaningful explanations
- Undermines trust and accountability

Governance Implications:

- Explainability and interpretability requirements
- Need for model documentation (model cards, fact sheets)
- Rights to explanation for affected individuals

Narahari

- Balance between accuracy and interpretability

4.3 Autonomy

Definition: AI systems can operate with varying degrees of independence from human control, making decisions and taking actions without direct human intervention.

Levels of Autonomy:

- **Low autonomy**: Human-in-the-loop (requires human approval for actions)
- **Medium autonomy**: Human-on-the-loop (human monitors but doesn't approve each action)
- **High autonomy**: Human-out-of-the-loop (system operates independently)

Why It Requires Governance:

- Reduced human oversight increases risk of unintended consequences
- Questions about accountability when autonomous systems cause harm
- Potential for systems to behave unexpectedly in novel situations
- Balance needed between efficiency and human control

Governance Implications:

- Requirements for human oversight appropriate to risk level
- "Kill switch" capabilities for high-risk systems
- Clear assignment of accountability
- Monitoring and intervention protocols

4.4 Speed and Scale

Definition: AI systems can process vast amounts of data and make millions of decisions at speeds far exceeding human capability.

Why It Requires Governance:

- Errors can proliferate rapidly and affect millions before detection
- Traditional human oversight mechanisms may be too slow
- Scale of impact magnifies consequences of bias or errors
- Difficult to halt or reverse harm once it occurs

Examples:

- Algorithmic trading systems executing millions of transactions per second
- Content moderation AI reviewing billions of social media posts

Narahari

- Credit scoring systems affecting millions of loan decisions

Governance Implications:

- Need for automated monitoring and alert systems
- Rapid incident response capabilities
- Pre-deployment testing at scale
- Continuous monitoring post-deployment

4.5 Data Dependency

Definition: AI systems rely heavily on data for training, validation, and operation. The quality, representativeness, and integrity of data directly affect system performance.

Why It Requires Governance:

- "Garbage in, garbage out"—poor data leads to poor AI
- Historical bias in data perpetuates discrimination
- Privacy concerns with personal data collection and use
- Data quality issues undermine reliability

Key Data Challenges:

- **Bias**: Data may not represent all populations equally
- **Quality**: Incomplete, inaccurate, or outdated data
- **Privacy**: Personal/sensitive information in training data
- **Provenance**: Unclear origins or ownership of data

Governance Implications:

- Data governance frameworks and policies
- Data quality standards and audits
- Privacy-enhancing technologies (PETs)
- Data minimization principles
- Transparency about data sources and characteristics

4.6 Probabilistic Outputs

Definition: AI systems typically provide probabilistic predictions rather than deterministic answers, meaning outputs come with uncertainty and potential for error.

Why It Requires Governance:

Narahari

- No system is 100% accurate—errors are inevitable
- Users may not understand or account for uncertainty
- Confidence levels may be miscommunicated or misunderstood
- Acceptable error rates vary by use case and context

Examples:

- Medical diagnosis AI might be 95% confident but still wrong
- Facial recognition with 98% accuracy still misidentifies 2% of people
- Credit scoring models predict probability of default, not certainty

Governance Implications:

- Clear communication of confidence levels and limitations
- Context-appropriate accuracy thresholds
- Human review for high-stakes decisions
- Understanding of error types (false positives vs. false negatives)
- Continuous monitoring of accuracy and performance

---

5. MODERN DRIVERS OF AI

5.1 Cloud Computing

What It Is: On-demand access to computing resources (servers, storage, processing power) via the internet without direct management by users.

How It Enables AI:

- Provides scalable computing power for training complex models
- Reduces upfront infrastructure investment
- Enables rapid experimentation and deployment
- Facilitates access to powerful GPUs and TPUs

Governance Considerations:

- Data security and privacy in cloud environments
- Vendor lock-in and dependency
- Compliance with data localization requirements
- Third-party risk management

5.2 Internet of Things (IoT)

Narahari

What It Is: Network of physical devices embedded with sensors, software, and connectivity that collect and exchange data.

How It Enables AI:

- Generates massive amounts of real-time data for AI systems
- Enables AI at the edge (processing on devices)
- Creates new applications (smart homes, autonomous vehicles, industrial IoT)
- Provides continuous feedback loops for AI improvement

Examples:

- Smart home devices learning user preferences
- Wearable health monitors providing data for AI diagnostics
- Industrial sensors enabling predictive maintenance
- Autonomous vehicle sensors feeding perception systems

Governance Considerations:

- Privacy concerns from pervasive data collection
- Security vulnerabilities in connected devices
- Data ownership and consent
- Edge AI governance challenges

5.3 Big Data

What It Is: Extremely large, complex datasets that traditional data processing methods cannot handle effectively, characterized by the "Five V's":

The Five V's of Big Data:

1. **Volume**: Enormous amounts of data
    - Terabytes to petabytes and beyond
    - From social media, sensors, transactions, etc.
2. **Velocity**: High speed of data generation and processing
    - Real-time or near-real-time data streams
    - Continuous data flow requiring rapid processing
3. **Variety**: Diverse data types and sources
    - Structured (databases), semi-structured (XML, JSON), unstructured (text, images, video)
    - Multiple sources requiring integration

Narahari

4. **Veracity**: Quality and accuracy of data
    - Trustworthiness and reliability
    - Noise, inconsistencies, and incomplete data
5. **Value**: Actionable insights extracted from data
    - Transforming data into business value
    - Identifying meaningful patterns and relationships

How It Enables AI:

- Provides the massive datasets needed to train sophisticated models
- Enables pattern recognition across diverse data sources
- Supports continuous learning and improvement
- Facilitates more accurate predictions through larger sample sizes

Governance Considerations:

- Data quality assurance and validation
- Privacy and consent at scale
- Data minimization vs. AI performance tradeoffs
- Bias amplification in large datasets

5.4 Advanced Algorithms and Computing Power

Key Developments:

- More efficient algorithms (e.g., transformer architectures)
- Specialized hardware (GPUs, TPUs, neuromorphic chips)
- Parallel processing capabilities
- Improved optimization techniques

Impact:

- Enables training of larger, more complex models
- Reduces time from months to days/hours for model training
- Makes AI more accessible and cost-effective
- Enables real-time AI applications

---

6. TYPES OF AI BASED ON CAPABILITY

6.1 Artificial Narrow Intelligence (ANI)

Narahari

Definition: AI designed to perform specific, well-defined tasks. Also called "Weak AI."

Characteristics:

- Excels at single, specific tasks
- Cannot generalize beyond its trained domain
- Most AI systems today are ANI
- Operates within predefined parameters

Examples:

- Virtual assistants (Siri, Alexa) for voice commands
- Image recognition systems
- Spam filters
- Recommendation engines (Netflix, Amazon)
- Chess or Go playing programs
- Language translation tools

Governance Focus:

- Task-specific risk assessment
- Clear definition of use case boundaries
- Monitoring for use beyond intended scope
- Domain-specific accuracy and fairness metrics

6.2 Artificial General Intelligence (AGI)

Definition: AI with human-like ability to understand, learn, and apply knowledge across a wide range of tasks. Also called "Strong AI."

Characteristics:

- Can perform any intellectual task a human can
- Transfers learning from one domain to another
- Understands context and adapts to new situations
- Exhibits reasoning and problem-solving across domains

Current Status:

- **Does NOT currently exist**
- Subject of ongoing research
- Timeline for achievement is uncertain and debated

Narahari

- Significant technical and theoretical challenges remain

Governance Considerations:

- Speculative but important for long-term planning
- Questions about rights, personhood, and accountability
- Need for adaptive governance frameworks
- International coordination on development

6.3 Artificial Super Intelligence (ASI)

Definition: Hypothetical AI that surpasses human intelligence across all domains, including creativity, social skills, and general wisdom.

Characteristics:

- Exceeds best human brains in every field
- Self-improving capabilities
- Potentially exponential intelligence growth
- Could fundamentally transform society

Current Status:

- **Purely theoretical**
- Subject of philosophical and ethical debate
- Focus of existential risk discussions
- Far beyond current technological capabilities

Governance Considerations:

- Long-term existential risk planning
- International cooperation frameworks
- Alignment problem (ensuring ASI goals align with human values)
- Precautionary principles in AGI research

---

7. AI USE CASES AND APPLICATIONS

7.1 Recognition/Perception

Definition: AI systems that identify, classify, or understand patterns in data (images, speech, text, etc.).

Types:

Narahari

Image Recognition:

- Facial recognition for security/authentication
- Medical image analysis (X-rays, MRIs, CT scans)
- Object detection in autonomous vehicles
- Quality control in manufacturing

Speech Recognition:

- Voice assistants and smart speakers
- Transcription services
- Voice-controlled systems
- Accessibility tools

Natural Language Processing (NLP):

- Sentiment analysis
- Text classification
- Named entity recognition
- Document analysis

Pattern Recognition:

- Anomaly detection in cybersecurity
- Fraud detection in financial transactions
- Biometric identification

Challenges:

- Bias in facial recognition (particularly affecting minorities)
- Privacy concerns with surveillance
- Accuracy disparities across demographic groups
- Misidentification risks and consequences

7.2 Forecasting/Prediction

Definition: AI systems that analyze historical data to predict future events, trends, or outcomes.

Applications:

Financial Services:

- Credit scoring and loan approvals
- Fraud detection

Narahari

- Investment recommendations
- Risk assessment

Healthcare:

- Disease diagnosis and prognosis
- Patient outcome predictions
- Hospital readmission risk
- Drug discovery and development

Business Operations:

- Demand forecasting
- Sales predictions
- Customer churn prediction
- Supply chain optimization

Public Sector:

- Predictive policing (controversial)
- Traffic management
- Weather forecasting
- Resource allocation

Challenges:

- Self-fulfilling prophecies (predictions influencing outcomes)
- Perpetuating historical biases
- Over-reliance on predictions with inherent uncertainty
- Lack of explainability in complex models
- Fairness issues in high-stakes decisions

7.3 Personalization/Recommendation

Definition: AI systems that tailor content, products, or services to individual users based on their preferences, behavior, and characteristics.

Applications:

Content Platforms:

- Video/music recommendations (Netflix, Spotify)
- News feed curation (social media)

Narahari

- Search result personalization
- Advertising targeting

E-Commerce:

- Product recommendations
- Dynamic pricing
- Personalized shopping experiences
- Email marketing customization

Education:

- Adaptive learning platforms
- Personalized curricula
- Student support systems

Healthcare:

- Personalized treatment plans
- Customized health recommendations
- Tailored medication dosages

Challenges:

- Filter bubbles and echo chambers
- Privacy concerns with behavioral tracking
- Manipulation and addiction through engagement optimization
- Lack of transparency in recommendation logic
- Discrimination in personalized pricing or access
- User autonomy and informed choice

---

8. AI GOVERNANCE FUNDAMENTALS

8.1 What is AI Governance?

Definition: AI governance encompasses the frameworks, policies, processes, and practices that organizations implement to ensure AI systems are developed and deployed responsibly, ethically, legally, and in alignment with organizational values.

Purpose:

- Manage risks associated with AI systems
- Ensure compliance with laws and regulations

Narahari

- Build trust with stakeholders
- Align AI use with organizational values and ethics
- Protect individuals and society from harm
- Enable responsible innovation

Scope:

- Technical safeguards and controls
- Organizational policies and procedures
- Risk management frameworks
- Ethical principles and guidelines
- Legal compliance mechanisms
- Stakeholder engagement processes

8.2 Why AI Governance is Needed

Risk Management:

- Preventing bias and discrimination
- Ensuring privacy and data protection
- Managing safety and security risks
- Addressing accountability gaps

Legal Compliance:

- Meeting regulatory requirements (EU AI Act, GDPR, etc.)
- Avoiding penalties and sanctions
- Demonstrating due diligence

Trust and Reputation:

- Building stakeholder confidence
- Maintaining brand reputation
- Demonstrating social responsibility

Operational Effectiveness:

- Improving AI system quality and reliability
- Reducing costly errors and failures
- Enabling sustainable AI programs

Narahari

8.3 Key AI Governance Principles

Multiple frameworks exist, but common themes emerge:

8.4 OECD AI Principles (2019)

Overview: Adopted by 42 countries, these principles represent international consensus on responsible AI.

The Five Principles:

1. Inclusive Growth, Sustainable Development, and Well-Being

- AI should benefit all people and the planet
- Enhance capabilities and improve quality of life
- Support sustainable development goals
- Address global challenges

2. Human-Centered Values and Fairness

- Respect for human rights and dignity
- Fair treatment and non-discrimination
- Diversity, inclusion, and equity
- Protection of vulnerable populations

3. Transparency and Explainability

- Meaningful information about AI systems
- Disclosure of AI use
- Explainability of decisions
- Access to information for affected individuals

4. Robustness, Security, and Safety

- Reliable and safe operation throughout lifecycle
- Security against attacks and threats
- Fallback plans and fail-safes
- Continuous monitoring and improvement

5. Accountability

- Clear assignment of responsibility
- Mechanisms for redress
- Ability to challenge decisions

Narahari

- Organizational accountability for impacts

Implementation Recommendations:

- Investment in AI R&D
- Fostering innovation ecosystems
- Enabling policy environments
- International cooperation
- Building AI literacy

8.5 Fair Information Practices (FIPs)

Historical Context: Originated in 1973 with the U.S. Department of Health, Education, and Welfare report; foundational to modern privacy laws.

Core Principles:

1. Notice/Awareness

- Inform individuals about data collection and use
- Clear, accessible privacy policies
- Transparent communication

2. Choice/Consent

- Individuals should have control over their data
- Meaningful consent for collection and use
- Options to opt-in or opt-out

3. Access/Participation

- Right to view personal data held about them
- Ability to correct inaccuracies
- Verification of data accuracy

4. Integrity/Security

- Data must be accurate and secure
- Protection against unauthorized access
- Regular security assessments

5. Enforcement/Redress

- Mechanisms to enforce principles
- Ability to seek remedies for violations

Narahari

- Accountability for non-compliance

Relevance to AI:

- Foundation for data governance in AI systems
- Informs requirements under GDPR, CCPA, and other privacy laws
- Essential for building trustworthy AI
- Addresses data dependency characteristic of AI

8.6 UNESCO Recommendation on the Ethics of AI

Adoption: Approved by 193 countries in November 2021

Key Values:

- Human rights and dignity
- Environmental and ecosystem flourishing
- Diversity and inclusiveness
- Peace and interconnectedness

Key Principles:

- Proportionality and Do No Harm
- Safety and security
- Right to privacy and data protection
- Multi-stakeholder governance
- Responsibility and accountability
- Transparency and explainability
- Human oversight and determination
- Sustainability
- Awareness and literacy
- Fairness and non-discrimination

Unique Contributions:

- Environmental considerations
- Cultural diversity emphasis
- Global South perspectives
- Holistic societal impact view

Narahari

9. AI GOVERNANCE FRAMEWORKS

9.1 NIST AI Risk Management Framework (AI RMF)

Purpose: Voluntary framework to help organizations manage risks associated with AI systems.

Development: Released in January 2023 by the U.S. National Institute of Standards and Technology.

Structure: Built around **Four Core Functions**:

1. GOVERN

- Establish policies, procedures, and culture for managing AI risks
- Define roles and responsibilities
- Engage stakeholders
- Ensure leadership commitment

2. MAP

- Understand context and identify risks
- Categorize AI system and assess intended use
- Analyze potential impacts
- Document assumptions and limitations

3. MEASURE

- Assess and benchmark AI system performance
- Evaluate against metrics for trustworthiness
- Conduct testing and validation
- Track and analyze metrics over time

4. MANAGE

- Allocate resources to mapped and measured risks
- Implement risk treatment strategies
- Prioritize actions based on risk assessment
- Monitor and communicate risk management decisions

Seven Trustworthy AI Characteristics:

1. **Valid and Reliable**: Consistently performs as intended
2. **Safe**: Free from harmful behavior and unintended consequences
3. **Secure and Resilient**: Protected against threats and able to recover
4. **Accountable and Transparent**: Clear responsibility and decision visibility

Narahari

5. **Explainable and Interpretable**: Understandable decisions and operations
6. **Privacy-Enhanced**: Protects personal information
7. **Fair with Harmful Bias Managed**: Equitable treatment and bias mitigation

Implementation Tools:

- NIST AI RMF Playbook (practical guidance)
- ARIA Program (Assessing Risks and Impacts of AI)
- Crosswalks to other frameworks (ISO, OECD, etc.)

Benefits:

- Flexible and customizable to different contexts
- Risk-based approach aligns with organizational practices
- Internationally recognized and referenced
- Complements existing risk management processes

9.2 IEEE (Institute of Electrical and Electronics Engineers)

Key Contribution: Ethically Aligned Design principles and standards

Focus Areas:

- Human rights emphasis
- Well-being prioritization
- Data agency
- Effectiveness and transparency
- Accountability
- Awareness of misuse potential
- Competence of creators

Standards Development:

- IEEE 7000 series on AI ethics
- Technical standards for implementation
- Industry-specific guidance
- Best practices documentation

---

10. AI TECHNOLOGY STACK

Narahari

Understanding the AI technology stack is essential for governance because risks and responsibilities exist at each layer.

10.1 Stack Layers (Bottom to Top)

---

1. COMPUTE (Hardware)

Components:

- **CPUs (Central Processing Units)**: General-purpose processors
- **GPUs (Graphics Processing Units)**: Parallel processing for AI training
- **TPUs (Tensor Processing Units)**: Google's specialized AI chips
- **NPUs/AI Accelerators**: Purpose-built AI hardware

Governance Considerations:

- Energy consumption and environmental impact
- Access to computational resources (equity concerns)
- Hardware vulnerabilities and security
- Supply chain dependencies

---

2. STORAGE

Components:

- Databases (relational, NoSQL)
- Data lakes and data warehouses
- Cloud storage services
- Distributed file systems

Governance Considerations:

- Data security and encryption
- Retention policies and deletion
- Backup and disaster recovery
- Privacy compliance (data localization, access controls)

---

3. NETWORK

Components:

- Internet connectivity

Narahari

- Data transfer protocols
- APIs (Application Programming Interfaces)
- Edge computing infrastructure

Governance Considerations:

- Data transmission security
- Bandwidth and latency impacts
- Cross-border data flows
- Network reliability and redundancy

---

4. PLATFORMS

Components:

- Cloud platforms (AWS, Azure, Google Cloud)
- AI development platforms (TensorFlow, PyTorch)
- MLOps platforms
- AI-as-a-Service offerings

Governance Considerations:

- Vendor lock-in and dependencies
- Service level agreements (SLAs)
- Third-party risk management
- Data governance in multi-tenant environments

---

5. SOFTWARE

Components:

- Operating systems
- AI frameworks and libraries
- Development tools
- Application software

Governance Considerations:

- Software vulnerabilities and patching
- Open-source license compliance
- Code quality and testing

Narahari

- Documentation and maintenance

---

6. OBSERVABILITY (Monitoring and Logging)

Components:

- Logging systems
- Monitoring dashboards
- Alerting mechanisms
- Audit trails

Governance Considerations:

- Continuous performance monitoring
- Detecting model drift and degradation
- Incident detection and response
- Compliance auditing and reporting

Importance for Governance:

- Enables accountability through audit trails
- Supports continuous improvement
- Facilitates incident investigation
- Demonstrates compliance

---

11. TYPES OF AI AND MACHINE LEARNING MODELS

11.1 Classic vs. Generative AI

Classic AI (Discriminative AI):

- **Purpose**: Classification, prediction, recognition
- **Output**: Categories, labels, predictions based on input
- **Examples**: Spam filters, image classifiers, fraud detection
- **Training**: Learns boundaries between classes in data

Generative AI:

- **Purpose**: Create new content similar to training data
- **Output**: Novel text, images, audio, video, code
- **Examples**: ChatGPT, DALL-E, Midjourney, GitHub Copilot
- **Training**: Learns underlying patterns and distributions in data

Narahari

Governance Differences:

- Generative AI raises unique concerns: deepfakes, misinformation, copyright
- Output quality and hallucinations in generative systems
- Increased need for content moderation and watermarking
- Different risk profiles requiring tailored governance

11.2 Proprietary vs. Open Source

Proprietary AI:

- **Characteristics**: Closed source code, vendor-controlled, often API-based access
- **Examples**: GPT-4 (OpenAI), Claude (Anthropic), proprietary enterprise solutions
- **Advantages**: Professional support, continuous updates, potentially higher quality
- **Disadvantages**: Lack of transparency, vendor dependency, cost

Open Source AI:

- **Characteristics**: Publicly available code, community-developed, modifiable
- **Examples**: Llama (Meta), BLOOM, Stable Diffusion base models
- **Advantages**: Transparency, customizability, no vendor lock-in, often free
- **Disadvantages**: Variable quality, limited support, security risks if not maintained

Governance Considerations:

- Open source enables auditing and explainability but may enable malicious use
- Proprietary models may be more secure but less transparent
- Licensing and compliance differ significantly
- Risk allocation between provider and deployer varies

11.3 Language Models vs. Multimodal Models

Language Models (LLMs - Large Language Models):

- **Focus**: Text understanding and generation
- **Inputs/Outputs**: Primarily text
- **Examples**: GPT-3/4, BERT, PaLM
- **Applications**: Chatbots, translation, summarization, code generation

Multimodal Models:

- **Focus**: Multiple data types (text, images, audio, video)

Narahari

- **Inputs/Outputs**: Combinations of text, images, audio, video
- **Examples**: GPT-4 with vision, CLIP, Flamingo
- **Applications**: Image captioning, visual question answering, multimedia search

Governance Considerations:

- Multimodal models have more complex risk profiles
- Potential for misuse increases with capabilities (deepfakes)
- Harder to test across multiple modalities
- Require diverse expertise for evaluation

11.4 SLMs vs. LLMs

Small Language Models (SLMs):

- Fewer parameters (millions to low billions)
- Lower computational requirements
- Faster inference
- Can run on edge devices
- Domain-specific often

Large Language Models (LLMs):

- Massive parameters (billions to trillions)
- Require significant computational resources
- Broader capabilities and knowledge
- Better at complex reasoning and generalization
- Typically cloud-based

Governance Considerations:

- SLMs may be more controllable and auditable
- LLMs pose greater risks but also greater capabilities
- Resource requirements affect accessibility
- Deployment location affects privacy and control

---

12. MACHINE LEARNING TRAINING APPROACHES

12.1 Supervised Learning

Definition: Training with labeled data where each input has a corresponding correct output.

Narahari

How It Works:

- Algorithm learns mapping from inputs to outputs
- Trained on examples with known correct answers
- Predicts labels for new, unseen data

Examples:

- Email spam classification (spam vs. not spam)
- Medical diagnosis (disease present vs. absent)
- Credit scoring (approve vs. deny)
- Image classification (cat vs. dog)

Requirements:

- Large amounts of labeled training data
- Labels must be accurate and representative
- Significant human effort for labeling

Governance Considerations:

- Label quality directly affects fairness and accuracy
- Labeler bias can propagate into model
- Data representativeness is critical
- Expensive and time-consuming to create labeled datasets

12.2 Unsupervised Learning

Definition: Training on unlabeled data to discover hidden patterns, structures, or relationships.

How It Works:

- Algorithm finds patterns without predefined labels
- Discovers groupings, associations, or anomalies
- No "correct answer" provided during training

Common Techniques:

- **Clustering**: Grouping similar items (customer segmentation)
- **Dimensionality Reduction**: Simplifying complex data
- **Anomaly Detection**: Identifying unusual patterns (fraud detection)
- **Association Rules**: Finding relationships (market basket analysis)

Examples:

Narahari

- Customer segmentation for marketing
- Anomaly detection in network security
- Topic modeling in document analysis
- Feature learning for other tasks

Governance Considerations:

- Harder to validate since no ground truth
- Potential for discovering spurious patterns
- Interpretation requires domain expertise
- May reveal sensitive patterns (proxy discrimination)

12.3 Reinforcement Learning

Definition: Learning through trial and error by interacting with an environment and receiving rewards/penalties.

How It Works:

- Agent takes actions in an environment
- Receives rewards for good actions, penalties for bad
- Learns policy to maximize cumulative reward
- Balances exploration (trying new things) with exploitation (using known strategies)

Key Components:

- **Agent**: The learner/decision-maker
- **Environment**: The system being interacted with
- **State**: Current situation
- **Action**: Choices available to agent
- **Reward**: Feedback signal (positive or negative)

Examples:

- Game-playing AI (AlphaGo, chess engines)
- Robotics and autonomous systems
- Resource optimization
- Recommendation systems with user feedback

Governance Considerations:

- Reward hacking (finding unintended ways to maximize reward)

Narahari

- Safety during exploration phase
- Ensuring reward function aligns with intended goals
- Long-term vs. short-term reward tradeoffs

12.4 Semi-Supervised Learning

Definition: Combines small amounts of labeled data with large amounts of unlabeled data.

Why It's Used:

- Labeling data is expensive and time-consuming
- Unlabeled data is often abundant and cheap
- Can achieve good performance with less labeled data

Governance Considerations:

- Balance between labeled and unlabeled data quality
- Potential for unlabeled data to introduce bias
- Validation complexity

12.5 Deep Learning

Definition: Subset of machine learning using artificial neural networks with multiple layers (hence "deep") to learn hierarchical representations.

Characteristics:

- Inspired by biological neural networks
- Automatically learns features from raw data
- Requires large amounts of data and computational power
- Excels at complex pattern recognition

Architecture Types:

- **Convolutional Neural Networks (CNNs)**: Image and video processing
- **Recurrent Neural Networks (RNNs)**: Sequential data, time series
- **Transformers**: Language models, attention mechanisms
- **Generative Adversarial Networks (GANs)**: Generating realistic synthetic data

Applications:

- Computer vision
- Natural language processing

Narahari

- Speech recognition
- Autonomous vehicles

Governance Considerations:

- "Black box" nature—difficult to interpret
- Requires massive datasets (privacy concerns)
- Prone to overfitting without proper regularization
- Vulnerability to adversarial attacks

12.6 Generative AI and Transformer Models

Generative AI:

- Creates new content (text, images, audio, video)
- Examples: GPT series, DALL-E, Midjourney, Stable Diffusion
- Based on learning data distributions

Transformer Architecture:

- Revolutionary architecture introduced in 2017 ("Attention is All You Need")
- Uses attention mechanisms to process sequences
- Parallel processing enables training on massive datasets
- Foundation of modern LLMs

Key Innovations:

- **Self-attention**: Weighing importance of different parts of input
- **Positional encoding**: Understanding sequence order
- **Scalability**: Can be trained on enormous datasets

Examples of Transformer-Based Models:

- GPT (Generative Pre-trained Transformer) series
- BERT (Bidirectional Encoder Representations from Transformers)
- T5 (Text-to-Text Transfer Transformer)

Governance Considerations:

- Hallucinations (generating false information confidently)
- Copyright and intellectual property issues
- Misinformation and deepfakes
- Bias amplification

Narahari

- Energy consumption for training
- Dual-use concerns (beneficial and harmful uses)

---

13. KEY DEFINITIONS FOR MODULE 1

Artificial Intelligence (AI): Theory and development of computer systems able to perform tasks normally requiring human intelligence.

Turing Test: Test of a machine's ability to exhibit intelligent behavior indistinguishable from a human.

Socio-Technical System: System involving both technical components (hardware, software) and social components (people, organizations, processes).

Complexity: Characteristic of AI systems referring to intricate architectures and numerous parameters making them difficult to understand.

Opacity/Black Box: Characteristic where AI decision-making processes are not transparent or easily interpretable.

Autonomy: Degree to which AI systems operate independently without human intervention.

Speed and Scale: AI's ability to process vast amounts of data and make millions of decisions rapidly.

Data Dependency: AI systems' reliance on data quality, representativeness, and integrity for performance.

Probabilistic Outputs: AI predictions that come with uncertainty rather than deterministic answers.

ANI (Artificial Narrow Intelligence): AI designed for specific tasks; cannot generalize beyond trained domain. Most current AI.

AGI (Artificial General Intelligence): Hypothetical AI with human-like ability to learn and apply knowledge across domains. Does not currently exist.

ASI (Artificial Super Intelligence): Theoretical AI surpassing human intelligence in all domains. Purely speculative.

AI Governance: Frameworks, policies, processes, and practices ensuring AI is developed and deployed responsibly, ethically, and legally.

OECD AI Principles: International principles for responsible AI covering inclusive growth, human-centered values, transparency, robustness, and accountability.

Fair Information Practices (FIPs): Foundational privacy principles covering notice, choice, access, integrity, and enforcement.

NIST AI RMF: Voluntary framework with four functions (Govern, Map, Measure, Manage) and seven trustworthy characteristics.

Narahari

Supervised Learning: Training with labeled data where inputs have corresponding correct outputs.

Unsupervised Learning: Training on unlabeled data to discover patterns without predefined labels.

Reinforcement Learning: Learning through trial-and-error by receiving rewards/penalties from environment.

Deep Learning: Machine learning using multi-layered neural networks to learn hierarchical representations.

Generative AI: AI that creates new content (text, images, audio) similar to training data.

Transformer: Neural network architecture using attention mechanisms, foundation of modern LLMs.

LLM (Large Language Model): Large-scale neural network trained on text data for language understanding and generation.

Narahari