AIGP Module 3 Advanced Practice Test

Responsible AI Governance and Risk Management

30 Medium to Hard Questions (Exam Simulation)

Question 1

A startup with 50 employees wants to implement AI governance but has limited resources. According to Module 3, what is the MOST appropriate approach for this organization?

A. Create entirely new AI-specific offices with dedicated staff
B. Combine AI governance with existing privacy or legal oversight functions and leverage existing assessment tools adapted for AI
C. Wait until the company is larger before implementing any AI governance
D. Hire external consultants to handle all governance functions

Correct Answer: B

Explanation:

- **Why B is correct:** The training explicitly states: "Smaller companies may need to combine services with privacy or legal oversight functions and leverage existing screening or risk assessment tools to include new AI aspects." This recognizes resource constraints while ensuring governance exists. Small organizations can integrate AI considerations into existing compliance, privacy, or risk management structures rather than creating entirely new functions.
- **Why A is incorrect:** The training contrasts small and large companies: "larger companies may create new AI-specific offices, oversight, and detailed processes" but small companies should leverage existing functions. Creating new offices would be resource-prohibitive for a 50-person startup.
- **Why C is incorrect:** The training emphasizes that governance should be implemented regardless of size, tailored appropriately. Waiting creates risk exposure and may embed poor practices that are harder to fix later. All organizations using AI need some governance structure.
- **Why D is incorrect:** While external help can supplement, organizations should develop internal capacity and not completely outsource governance. The training recommends internal or internal/external hybrid structures with clear accountability, which requires internal involvement.

Question 2

According to the NIST AI Risk Management Framework, which function is described as being "infused throughout AI risk management and enables the other functions"?

Narahari

A. Map
B. Measure
C. Manage
D. Govern

Correct Answer: D

Explanation:

- **Why D is correct:** The training explicitly states: "Govern: Cultivate and implement a culture of risk management; this function is infused throughout AI risk management and enables the other functions." Govern establishes the foundation, culture, roles, and structures that make the other three functions (Map, Measure, Manage) possible and effective.
- **Why A is incorrect:** "Map: Identify use and risks related to use" is about identifying what AI systems exist and their risks, not the foundational function that enables others.
- **Why B is incorrect:** "Measure: Assess, analyze and track risks" focuses on quantifying and monitoring risks, not establishing the enabling governance structure.
- **Why C is incorrect:** "Manage: Prioritize risks and act based on projected impact" is about taking action on risks, not creating the foundational governance culture that enables all functions.

Question 3

An organization is deciding between centralized, decentralized, and hybrid AI governance models. A large multinational company with diverse business units wants consistency in core principles but needs flexibility for local implementation. Which model is MOST appropriate?

A. Centralized model for complete control
B. Decentralized model for maximum flexibility
C. Hybrid model combining central policy direction with local implementation authority
D. No formal governance structure

Correct Answer: C

Explanation:

- **Why C is correct:** The training describes the hybrid model as: "allows for a combination of centralized and local governance; typically seen when a large organization assigns an individual the main responsibility for AI affairs and local entities fulfill and support the policies and directives from the central governing body." This balances the need for consistent principles (centralized) with local flexibility (decentralized), which perfectly fits the scenario's requirements.
- **Why A is incorrect:** A purely centralized model where "one team or person responsible for AI-related affairs and all other people or organizations flow through this point" would create bottlenecks in a large, diverse multinational organization and lack the flexibility for local context.

Narahari

- **Why B is incorrect:** A fully decentralized model "involves delegating decision-making authority to the lower levels" but risks inconsistency in core principles across business units. The scenario explicitly wants consistency in core principles, which pure decentralization doesn't ensure.
- **Why D is incorrect:** The training emphasizes that governance structure is essential. Having no formal structure creates risk exposure, inconsistency, and accountability gaps.

Question 4

A risk assessment identifies an AI system for smartphone facial recognition with severity rating of 2/3 and probability of 2/3 on a 3x3 matrix. Another system for law enforcement facial recognition has severity 3/3 and probability 3/3. Using the risk scoring methodology described in the training, what are the risk scores?

A. Smartphone: 4, Law enforcement: 6
B. Smartphone: 4, Law enforcement: 9
C. Smartphone: 2, Law enforcement: 3
D. Both have equal risk scores

Correct Answer: B

Explanation:

- **Why B is correct:** The training states: "Risk scoring can be calculated by first rating the risk on its severity and probability of occurring; then, multiplying the severity of the harm with the probability of the occurrence." For smartphone: 2 × 2 = 4. For law enforcement: 3 × 3 = 9. The training specifically uses facial recognition as an example: "Facial recognition for an individual's smartphone is a relatively low risk, but facial recognition used by law enforcement on a wide-scale basis could be a relatively high risk."
- **Why A is incorrect:** This adds severity and probability (2+2=4, 3+3=6) instead of multiplying them. The training explicitly says to multiply, not add.
- **Why C is incorrect:** This uses only the severity ratings without applying the probability factor. The methodology requires considering both dimensions and multiplying them.
- **Why D is incorrect:** The scenarios have different severity and probability ratings, so they cannot have equal risk scores. The law enforcement application has higher stakes (severity) and broader deployment (probability), yielding a higher score.

Question 5

According to the NIST AI RMF, which characteristic of trustworthy AI specifically addresses the need for AI systems to withstand attacks and continue functioning appropriately even under adverse conditions?

A. Valid and reliable
B. Safe

Narahari

C. Secure and resilient
D. Accountable and transparent

Correct Answer: C

Explanation:

- **Why C is correct:** "Secure and resilient" specifically addresses protection against attacks and continued appropriate functioning under adverse conditions. Security covers protection from cyber threats, while resilience addresses the ability to recover and maintain functionality despite disruptions.
- **Why A is incorrect:** "Valid and reliable" means the AI performs its intended function accurately and consistently, not that it withstands attacks or adverse conditions.
- **Why B is incorrect:** "Safe" addresses preventing physical harm and unacceptable safety risks, not specifically about withstanding attacks or maintaining function under adverse conditions.
- **Why D is incorrect:** "Accountable and transparent" addresses organizational responsibility and openness about how AI works, not technical resilience against attacks.

Question 6

An organization is performing a risk assessment for an AI system. According to the training, which of the following is NOT listed as a key criterion for risk assessment?

A. Business purpose and planned uses of the AI
B. Potential harms, including false positives and negative predictions
C. The marketing budget allocated to promote the AI system
D. Third-party risks from involving external vendors

Correct Answer: C

Explanation:

- **Why C is correct:** Marketing budget is not mentioned as a risk assessment criterion. The training lists specific criteria: "Business purpose and planned uses of the AI, Potential harms including false positives and negative predictions, Descriptions of the data used to train the AI including sensitive data, Functionality, Performance metrics, Benchmarking, Third-party risks." Marketing budget is a business consideration but not a risk assessment criterion.
- **Why A is incorrect:** This is explicitly listed: "Business purpose and planned uses of the AI" is a key risk assessment criterion to understand what the AI is intended to do and how it will be used.
- **Why B is incorrect:** This is explicitly listed: "Potential harms, including false positives and negative predictions" is a critical risk criterion for understanding possible impacts.

Narahari

- **Why D is incorrect:** This is explicitly listed: "Third-party risks: Determine and include the risks raised by involving third parties" is a key criterion since many organizations use external AI vendors or services.

Question 7

According to Module 3, what is the PRIMARY purpose of maintaining an AI inventory with risk scores attached to each system?

A. To comply with marketing requirements
B. To allocate appropriate resources and determine review frequency based on risk levels
C. To increase the number of AI systems deployed
D. To eliminate all AI-related risks

Correct Answer: B

Explanation:

- **Why B is correct:** The training explicitly states: "One way to do this is to inventory all your AI systems and attach a risk score to each system. This risk score will help you allocate appropriate resources to that system. It will also highlight the frequency with which you need to review the algorithm to evaluate whether it still meets its purpose." The inventory enables risk-based resource allocation and appropriate review scheduling.
- **Why A is incorrect:** Marketing is not mentioned as a purpose for AI inventory. The inventory serves governance, risk management, and monitoring purposes, not marketing functions.
- **Why C is incorrect:** The purpose isn't to increase AI deployment—it's to manage existing systems appropriately. More systems isn't the goal; proper governance of systems is the goal.
- **Why D is incorrect:** Risk cannot be completely eliminated. The inventory helps identify, prioritize, and manage risks appropriately, not eliminate them entirely. The training emphasizes risk management, not risk elimination.

Question 8

An AI governance team is deciding who should be involved in establishing governance. According to the training, which statement BEST reflects the recommended approach?

A. Only technical teams should be involved since they understand the technology
B. Cross-functional collaboration including privacy, legal, ethics, technical teams, and business stakeholders ensures alignment of policy goals, business goals, and technical realities
C. Governance should be handled exclusively by external consultants
D. Only executive leadership needs to be involved

Correct Answer: B

Explanation:

Narahari

- **Why B is correct:** The training explicitly states: "Designers, developers, builders, marketers and managers should also be part of AI governance. This provides an opportunity to communicate and helps ensure policy goals, business goals and tech realities align." It also emphasizes establishing "cross-functional collaboration in the AI governance program" and lists numerous roles from privacy officers to data scientists to marketing representatives.
- **Why A is incorrect:** While technical expertise is important, the training emphasizes that governance requires multiple perspectives. "Privacy, ethics, RAI (responsible AI), governance and legal personnel all look at similar aspects of the same things" and should work together with technical teams.
- **Why C is incorrect:** The training recommends internal governance structures with clear roles and responsibilities. External consultants can supplement but shouldn't exclusively handle governance, as the organization must maintain accountability.
- **Why D is incorrect:** While leadership support is crucial ("AI governance leadership support at the earliest opportunity is crucial"), the training emphasizes involving multiple stakeholder groups, not limiting involvement to executives only.

Question 9

According to the training, risk assessment for AI systems is described as "context-specific." Which of the following factors determine this context? (Select the MOST complete answer)

A. Owner/operator, specific industry/use case, potential social impacts, timing, and jurisdictional controls
B. Only the technical complexity of the AI model
C. Only the size of the organization
D. Only the cost of the AI system

Correct Answer: A

Explanation:

- **Why A is correct:** The training explicitly states: "Risk assessment is critical to the successful governance of AI systems, but is context-specific as to the: Owner and operator, Specific industry and use case, Potential social impacts, Timing and use of AI, Jurisdictional controls." All these factors must be considered to properly assess risk in context.
- **Why B is incorrect:** Technical complexity is one consideration but far from complete. Context includes who uses it, where, when, and for what purpose, not just technical characteristics.
- **Why C is incorrect:** While organization size affects governance approach, risk assessment context requires considering the specific use case, impacts, jurisdiction, and other factors beyond just organizational size.

Narahari

- **Why D is incorrect:** Cost is a business consideration but doesn't determine risk context. A low-cost system can be high-risk (e.g., used in critical decisions), and an expensive system can be low-risk.

Question 10

Model cards or fact sheets are mentioned as critical documentation components. According to the training, what information should these documents provide?

A. Only the technical architecture of the model
B. Standardized information about the model including its function, output, intended use, and limitations
C. Only the names of the developers
D. Only the cost to develop the model

Correct Answer: B

Explanation:

- **Why B is correct:** The training states: "Model cards or facts sheets - Provide standardized information about the model and its function/output." The purpose is to communicate essential information about what the model does, how it works, what it's intended for, and what its limitations are in a standardized format. This supports transparency, proper use, and audit requirements.
- **Why A is incorrect:** While technical architecture might be included, model cards provide broader information including function, output, intended use, and limitations—not just technical architecture.
- **Why C is incorrect:** Developer names might be included for accountability, but the primary purpose is documenting the model's function, capabilities, and limitations, not personnel information.
- **Why D is incorrect:** Development cost is a business consideration that might be tracked separately, but it's not the focus of model cards, which document the model's purpose, function, performance, and limitations.

Question 11

A large financial services company is implementing AI governance. According to the training, what is a recommended early step for gaining leadership support?

A. Implement governance without informing leadership
B. Identify leadership already using AI who would support improved governance, explain how responsible AI is a differentiator, and show how the organization can govern AI
C. Focus only on technical implementation without business justification
D. Wait until an AI incident occurs before engaging leadership

Correct Answer: B

Explanation:

Narahari

- **Why B is correct:** The training provides "Three considerations for engaging leadership when kicking off an AI governance program: 1. Identify any leadership already using AI who would support improved governance and structure... 2. Explain how responsible AI is a differentiator... 3. Show why and how your organization and leadership can govern AI." This strategic approach builds support by connecting to existing initiatives, demonstrating business value, and showing feasibility.
- **Why A is incorrect:** The training explicitly states "AI governance leadership support at the earliest opportunity is crucial" and "A holistic organizational approach with support from leadership can help influence behavioral and cultural change." Implementing without leadership engagement undermines success.
- **Why C is incorrect:** Business justification and leadership support are essential. The training emphasizes "Define the business case" and connecting governance to organizational goals, not just technical implementation.
- **Why D is incorrect:** Waiting for an incident is reactive and exposes the organization to unnecessary risk. The training advocates proactive governance with early leadership engagement, not waiting until problems occur.

Question 12

According to Module 3, what is the relationship between different organizational risk management strategies (security, privacy, business) regarding AI?

A. They should be kept completely separate with no interaction
B. All risk management strategies need to be aligned and integrated; if they don't intersect, gaps may be exploited and surface through incidents
C. Only technical security risk management matters for AI
D. Risk management strategies are optional for AI systems

Correct Answer: B

Explanation:

- **Why B is correct:** The training explicitly states: "All of an organization's risk management strategies need to be aligned. If they do not intersect, there will be gaps between them which may be exploited and surface through incidents." This emphasizes the need for integrated or holistic risk management for AI that brings together security, privacy, business, and other risk perspectives.
- **Why A is incorrect:** This contradicts the training's emphasis on alignment and integration. Siloed risk management creates the gaps that the training warns against.
- **Why C is incorrect:** The training discusses multiple risk types: "Strategies: security/operational risk, privacy risk, business risk (integrated or holistic for AI)." All are important, not just technical security.
- **Why D is incorrect:** The training emphasizes risk management as essential, not optional. The entire module focuses on establishing comprehensive risk management for AI systems.

Narahari

Question 13

An organization discovers that its deployed AI system's algorithm has drifted from its original purpose. According to the training, what should the AI inventory and monitoring process enable the organization to do?

A. Ignore the drift since the system is already deployed
B. Evaluate whether the system still meets its purpose, assess drift, and determine how to allocate auditing resources
C. Immediately deploy more AI systems
D. Only check the system once at initial deployment

Correct Answer: B

Explanation:

- **Why B is correct:** The training states: "It will also highlight the frequency with which you need to review the algorithm to evaluate whether it still meets its purpose. You should also evaluate if there has been drift or changes in the algorithm, as well as how to allocate auditing resources." The inventory and monitoring process is designed to catch exactly this type of issue and enable appropriate response.
- **Why A is incorrect:** Drift is a serious governance concern that must be addressed. The training emphasizes continuous monitoring and evaluation, not ignoring changes in deployed systems.
- **Why C is incorrect:** Discovering drift in an existing system doesn't call for deploying more systems—it requires addressing the drift in the current system through evaluation, potential retraining, or deactivation.
- **Why D is incorrect:** The training emphasizes ongoing monitoring: "Continuously improve the system by retraining with new data as needed and with human input and feedback." One-time checking at deployment is insufficient.

Question 14

According to the NIST AI RMF's seven characteristics of trustworthy AI, which characteristic specifically addresses the need for AI systems to protect individuals' data and implement appropriate privacy safeguards?

A. Valid and reliable
B. Explainable and interpretable
C. Privacy-enhanced
D. Accountable and transparent

Correct Answer: C

Explanation:

- **Why C is correct:** "Privacy-enhanced" is the specific NIST AI RMF characteristic addressing privacy protections and data safeguards. This characteristic requires

Narahari

implementing privacy-enhancing technologies and practices to protect personal information.

- **Why A is incorrect:** "Valid and reliable" addresses whether the AI performs its intended function accurately and consistently, not privacy protections.
- **Why B is incorrect:** "Explainable and interpretable" addresses understanding how AI makes decisions, not specifically privacy and data protection.
- **Why D is incorrect:** "Accountable and transparent" addresses organizational responsibility and openness about AI operations, which relates to but isn't specifically about privacy protections and data safeguards.

Question 15

An organization is determining roles for AI governance. According to the training, what value do non-AI engineers bring to AI governance?

A. No value since they don't understand AI
B. They help generate questions about AI generally and provide a strong operations perspective critical for implementation and product release processes
C. They should only handle non-AI systems
D. They replace AI engineers in governance decisions

Correct Answer: B

Explanation:

- **Why B is correct:** The training explicitly states: "Non-AI engineers will: help generate questions about AI generally; have a strong operations perspective that will be critical to determine how to take AI capabilities and drive them into implementation and product release processes." Non-AI engineers provide valuable operational and implementation perspectives that AI specialists might lack.
- **Why A is incorrect:** The training explicitly identifies valuable contributions from non-AI engineers, recognizing that AI governance requires diverse perspectives, not just AI expertise.
- **Why C is incorrect:** Non-AI engineers should be included in AI governance, not excluded from it. Their operational perspective contributes to comprehensive governance.
- **Why D is incorrect:** Non-AI engineers complement, not replace, AI engineers. The training emphasizes cross-functional collaboration with both AI and non-AI engineers contributing their respective expertise.

Question 16

According to the training, when should impact assessments be performed in the AI lifecycle?

Narahari

A. Only after deployment when problems occur
B. Only at the initial design phase
C. Throughout the AI lifecycle, starting in the model design phase
D. Impact assessments are not necessary for AI systems

Correct Answer: C

Explanation:

- **Why C is correct:** The training states: "In the AI model design phase, it is important to perform or review an impact assessment on the AI model. An impact assessment is a risk management tool used to assess an AI system's benefits, risks and limitations throughout its life cycle." The emphasis on "throughout its life cycle" indicates ongoing assessment, not just at one phase.
- **Why A is incorrect:** Waiting until after deployment when problems occur is reactive. The training emphasizes proactive risk assessment starting early in the design phase and continuing throughout the lifecycle.
- **Why B is incorrect:** While impact assessments should start early (design phase), they shouldn't stop there. The training emphasizes assessment "throughout its life cycle," not just initially.
- **Why D is incorrect:** Impact assessments are explicitly recommended: "it is important to perform or review an impact assessment." They're a key risk management tool for AI governance.

Question 17

An organization is building documentation for its AI systems. According to the training, why is it critical to document all decisions the governance group makes, even those not addressing regulatory requirements?

A. Documentation is only needed for regulatory compliance
B. Documentation is critical for informing future audits and maintaining institutional knowledge
C. Documentation has no value if not required by regulation
D. Only technical documentation matters

Correct Answer: B

Explanation:

- **Why B is correct:** The training states: "All decisions the group makes should be documented, whether the decisions address regulatory requirements or not" and "Document all decisions and updates — these will be critical for informing future audits." Documentation serves multiple purposes beyond compliance: institutional knowledge preservation, audit support, accountability, and enabling continuous improvement.
- **Why A is incorrect:** The training explicitly says documentation should include decisions "whether the decisions address regulatory requirements or not," indicating documentation serves purposes beyond just regulatory compliance.

Narahari

- **Why C is incorrect:** The training identifies multiple values for documentation: audit preparation, institutional knowledge, accountability, and continuous improvement—not just regulatory compliance.
- **Why D is incorrect:** The training recommends documenting "model planning, design and development steps" and "all decisions the group makes," which includes governance and business decisions, not just technical documentation.

Question 18

According to the training, what is the recommended approach for organizations that are small or under-resourced regarding AI documentation and auditing?

A. Skip documentation entirely
B. Many companies and resources are available to help with AI auditing and documentation
C. Wait until the organization is larger
D. Only document what's legally required

Correct Answer: B

Explanation:

- **Why B is correct:** The training explicitly states: "If your organization is small or under-resourced, many companies and resources are available to help with AI auditing and documentation." This acknowledges resource constraints while emphasizing that documentation and auditing remain necessary, with external assistance available.
- **Why A is incorrect:** Documentation is critical regardless of organization size. The training recommends seeking external help for small organizations, not skipping documentation.
- **Why C is incorrect:** Small organizations still need governance and documentation. The training recommends adapting approaches to size constraints and using available resources, not waiting to grow larger.
- **Why D is incorrect:** The training emphasizes documenting all decisions "whether the decisions address regulatory requirements or not," not limiting documentation to legal requirements.

Question 19

According to Module 3, what should be included in documentation about counterfactual explanations?

A. Only the final model output
B. Details on what new or different input may affect the output of the AI process
C. Only the training data sources
D. Only the model's accuracy metrics

Correct Answer: B

Narahari

Explanation:

- **Why B is correct:** The training specifically lists "Counterfactual explanations - Details on what new or different input may affect the output of the AI process" as part of required documentation. Counterfactual explanations help stakeholders understand how changing inputs would change outputs, supporting explainability and contestability.
- **Why A is incorrect:** Counterfactual explanations aren't about final output—they're about how different inputs would produce different outputs, helping users understand the decision logic.
- **Why C is incorrect:** While training data sources should be documented separately, counterfactual explanations specifically address how input variations affect outputs, not training data sources.
- **Why D is incorrect:** Accuracy metrics are important documentation but are separate from counterfactual explanations, which focus on input-output relationships and decision boundaries.

Question 20

An organization wants to build on existing processes for AI impact assessments. According to the training, what existing assessments can be leveraged?

A. Only financial audits
B. Data protection impact assessments (DPIAs) or privacy impact assessments (PIAs)
C. Only marketing assessments
D. No existing assessments are relevant to AI

Correct Answer: B

Explanation:

- **Why B is correct:** The training explicitly states: "Utilize existing resources and processes • Where possible, build off existing data protection impact assessments (DPIAs) or privacy impact assessments (PIAs)." This allows organizations to leverage established assessment processes while adapting them for AI-specific considerations.
- **Why A is incorrect:** Financial audits serve different purposes and aren't mentioned as a foundation for AI impact assessments, which focus on risks, benefits, and limitations of AI systems.
- **Why C is incorrect:** Marketing assessments aren't mentioned as relevant foundations. The training specifically identifies DPIAs and PIAs as the existing assessments to build upon.
- **Why D is incorrect:** The training explicitly recommends building on existing DPIAs and PIAs, showing that existing assessment processes are relevant and can be adapted for AI.

Question 21

According to the NIST AI RMF, which function involves identifying AI uses and the risks related to those uses?

A. Govern
B. Map
C. Measure
D. Manage

Correct Answer: B

Explanation:

- **Why B is correct:** The training defines the Map function as: "Map: Identify use and risks related to use." This function is about discovering and cataloging what AI systems exist, how they're used, and what risks they present.
- **Why A is incorrect:** "Govern: Cultivate and implement a culture of risk management" is about establishing governance foundations, not identifying specific uses and risks.
- **Why C is incorrect:** "Measure: Assess, analyze and track risks" comes after identification (Map). You must first identify risks (Map) before you can assess and track them (Measure).
- **Why D is incorrect:** "Manage: Prioritize risks and act based on projected impact" focuses on taking action on risks, not identifying them initially.

Question 22

An AI system's risk assessment shows high severity (3/3) but low probability (1/3) of harm occurring. Using the methodology described in the training, what is the risk score, and what does this suggest about governance priorities?

A. Risk score: 4; this is a low priority
B. Risk score: 3; this requires moderate attention to severity despite low probability
C. Risk score: 3; high-severity risks may warrant attention even with lower probability
D. Risk score: 2; this can be ignored

Correct Answer: C

Explanation:

- **Why C is correct:** Using the formula "multiplying the severity of the harm with the probability of the occurrence," the risk score is 3 × 1 = 3. This represents a moderate overall risk score. However, the high severity (3/3) suggests that if the harm does occur, it would be catastrophic. Governance should consider both the quantitative score and the fact that high-severity risks may require attention even when probability is lower, particularly if mitigation can reduce probability further or severity impacts.
- **Why A is incorrect:** The risk score is 3 (3 × 1), not 4. Also, high-severity risks shouldn't automatically be considered low priority even if probability is low, especially if mitigation is feasible.

Narahari

- **Why B is incorrect:** While the risk score calculation (3) is correct, the explanation should note that the high severity component means this could warrant significant attention despite lower probability, not just "moderate attention."
- **Why D is incorrect:** The risk score is 3, not 2. More importantly, high-severity risks should not be ignored even with lower probability—they may require preventive measures given the catastrophic potential.

Question 23

According to the training, what is an important consideration when engaging stakeholders in AI governance?

A. Only include technical experts
B. Aim for diversity of inputs by measures such as age, gender, race, region, and culture
C. Exclude end users from the process
D. Only involve stakeholders after deployment

Correct Answer: B

Explanation:

- **Why B is correct:** The training explicitly states: "Aim for a diversity of inputs by other measures for testing and assessment, such as age, gender, race, region, culture, etc." Diverse stakeholder engagement helps identify risks and impacts that homogeneous teams might miss.
- **Why A is incorrect:** The training emphasizes cross-functional engagement beyond just technical experts: "Privacy, ethics, RAI, governance and legal personnel" and "Include the users when possible."
- **Why C is incorrect:** The training explicitly recommends: "Include the users when possible." User input is valuable for understanding real-world impacts and usability.
- **Why D is incorrect:** The training emphasizes: "Involve stakeholders early in the process." Early engagement prevents problems and ensures governance aligns with needs, rather than trying to fix issues post-deployment.

Question 24

According to Module 3, what should organizations do with AI systems identified through inventory and monitoring as needing deactivation or localization?

A. Continue operating them without changes
B. Ensure there is a procedure in place to deactivate a system or localize it as needed
C. Immediately deploy replacement systems
D. Remove all documentation about the system

Correct Answer: B

Explanation:

Narahari

- **Why B is correct:** The training states: "Ensure there is a procedure in place to deactivate a system or localize it as needed." This recognizes that some AI systems may need to be shut down or restricted in scope when they're not meeting objectives, causing harm, or violating requirements.
- **Why A is incorrect:** If a system needs deactivation or localization, continuing without changes would perpetuate problems or risks. The training emphasizes monitoring and taking action when needed.
- **Why C is incorrect:** Deactivation doesn't automatically mean immediate replacement. The decision to deploy alternatives should be based on need and proper governance, not automatic replacement.
- **Why D is incorrect:** Documentation should be preserved even for deactivated systems for audit trails, lessons learned, and accountability. The training emphasizes comprehensive documentation, not destroying it.

Question 25

According to the training, algorithmic impact assessments should cover data issues and document what types of decisions?

A. Only technical architecture decisions
B. Only compliance-related decisions
C. Risk identification and mitigation decisions, and identifying who approves and accepts risk on behalf of the organization
D. Only budget allocation decisions

Correct Answer: C

Explanation:

- **Why C is correct:** The training states: "An algorithmic impact assessment should cover the data issues and document decisions your stakeholder group makes. This may include risk identification and mitigation or identifying who approves and accepts risk on behalf of your organization." This comprehensive approach ensures accountability and transparency in risk management.
- **Why A is incorrect:** Algorithmic impact assessments cover much more than technical architecture—they include risk identification, mitigation strategies, data issues, and governance decisions.
- **Why B is incorrect:** While compliance is important, the training indicates impact assessments should document all stakeholder decisions about risks and mitigation, not just compliance-related decisions.
- **Why D is incorrect:** Budget decisions aren't the focus of algorithmic impact assessments. These assessments focus on risks, harms, benefits, data issues, and risk acceptance decisions.

Narahari

Question 26

A governance team is defining the business case for AI use. According to the training, what questions should stakeholders address?

A. Only the potential revenue increase
B. Cost/benefit analysis, tradeoffs in using AI/ML vs. other solutions, and whether AI is the right solution for the mission and purpose
C. Only the technical feasibility
D. Only the timeline for implementation

Correct Answer: B

Explanation:

- **Why B is correct:** The training states: "Define the business case... What is the cost/benefit analysis? What are the tradeoffs in using AI/ML vs. other solutions? ... Ask the stakeholder group to assess whether AI is the right solution. Is it suitable for the mission and for the purpose that needs to be addressed?" This comprehensive evaluation ensures AI is used appropriately.
- **Why A is incorrect:** The business case involves more than revenue—it includes costs, risks, tradeoffs with alternatives, and mission suitability. A narrow revenue focus ignores important considerations.
- **Why C is incorrect:** While technical feasibility matters, the business case also requires evaluating cost/benefit, comparing alternatives, and assessing mission alignment—not just whether it's technically possible.
- **Why D is incorrect:** Timeline is one factor, but the training emphasizes comprehensive evaluation including cost/benefit, comparison with alternatives, and whether AI is appropriate for the purpose.

Question 27

According to the training, what role do researchers play in AI governance?

A. No role—only developers matter
B. Help identify key risks for the organization and core principles to uphold
C. Only handle marketing materials
D. Only manage budgets

Correct Answer: B

Explanation:

- **Why B is correct:** The training explicitly states: "Researchers can help identify key risks for the organization and core principles to uphold." Researchers bring analytical and investigative perspectives that help identify emerging risks and establish ethical foundations.
- **Why A is incorrect:** The training lists researchers alongside developers, data scientists, and AI/ML engineers as important contributors to governance, not as irrelevant.

Narahari

- **Why C is incorrect:** Marketing isn't mentioned as researchers' governance role. The training emphasizes their contribution to risk identification and principle development.
- **Why D is incorrect:** Budget management isn't mentioned as researchers' role. Their contribution is intellectual—identifying risks and establishing principles for governance.

Question 28

According to the NIST AI RMF, what does the "Manage" function involve?

A. Only documenting risks
B. Prioritize risks and act based on projected impact
C. Only identifying risks
D. Ignoring low-probability risks

Correct Answer: B

Explanation:

- **Why B is correct:** The training defines the Manage function as: "Manage: Prioritize risks and act based on projected impact." This function is about taking action—deciding which risks need attention first and implementing mitigation measures based on their potential impact.
- **Why A is incorrect:** Documenting is important but the Manage function is about active risk prioritization and taking action, not just documentation.
- **Why C is incorrect:** Identifying risks is the "Map" function. "Manage" comes later and involves deciding what to do about identified and measured risks.
- **Why D is incorrect:** Risk management involves prioritization based on impact, not ignoring categories of risk. Even low-probability risks might need attention if they have high severity.

Question 29

An organization discovers through monitoring that an AI system needs continuous improvement. According to the training, what should this involve?

A. Leaving the system unchanged
B. Retraining with new data as needed and incorporating human input and feedback
C. Immediately replacing the entire system
D. Removing all documentation

Correct Answer: B

Explanation:

- **Why B is correct:** The training states: "Continuously improve the system by retraining with new data as needed and with human input and feedback." This approach maintains system relevance and performance through iterative improvement based on monitoring insights.

Narahari

- **Why A is incorrect:** Continuous improvement by definition requires changes. Leaving the system unchanged contradicts the monitoring insights indicating improvement is needed.
- **Why C is incorrect:** Improvement doesn't necessarily mean complete replacement. The training emphasizes iterative enhancement through retraining and feedback, which is typically more practical than full replacement.
- **Why D is incorrect:** Documentation should be maintained and updated, not removed. The training emphasizes comprehensive documentation throughout the lifecycle.

Question 30

According to Module 3, when establishing AI governance, what is the value of leveraging existing structures and building on previous lessons?

A. It wastes time—everything should be built from scratch
B. It helps support build-out efficiently by using proven frameworks and institutional knowledge
C. Existing structures are never relevant to AI
D. Only new structures work for AI governance

Correct Answer: B

Explanation:

- **Why B is correct:** The training recommends: "Involve stakeholders early in the process... Leverage existing structures, and build on previous lessons as they apply to AI." This efficient approach uses institutional knowledge and proven frameworks, adapting them for AI-specific needs rather than starting from nothing.
- **Why A is incorrect:** Leveraging existing structures is explicitly recommended, not criticized as time-wasting. Starting from scratch ignores valuable institutional knowledge and proven approaches.
- **Why C is incorrect:** While AI has unique aspects, many existing governance structures (privacy, security, risk management) are relevant and can be adapted. The training recommends building on existing foundations.
- **Why D is incorrect:** The training explicitly recommends leveraging existing structures and lessons learned, recognizing that many governance principles and processes can be adapted rather than created anew.

End of Module 3 Advanced Practice Test

Scoring Guide:

- **27-30 correct:** Exceptional understanding - exam-ready
- **24-26 correct:** Strong understanding - minor review needed
- **21-23 correct:** Good foundation - review NIST AI RMF functions

Narahari

- **18-20 correct:** Adequate knowledge - focus on governance structures and risk assessment
- **Below 18:** Review Module 3 thoroughly, especially NIST AI RMF and risk management methodologies

Key Concepts to Master for Module 3:

1. **NIST AI RMF Four Functions:**
    - Govern, Map, Measure, Manage (know what each does)
2. **Seven Trustworthy AI Characteristics:**
    - Valid/reliable, safe, secure/resilient, accountable/transparent, explainable/interpretable, privacy-enhanced, fair/bias-mitigated
3. **Governance Models:**
    - Centralized, decentralized, hybrid (when to use each)
4. **Risk Assessment:**
    - Severity × Probability methodology
    - Context-specific factors
    - Key assessment criteria
5. **Roles and Responsibilities:**
    - Cross-functional team composition
    - What different roles contribute
6. **Documentation Requirements:**
    - Model cards, impact assessments, decision documentation
    - Why documentation matters beyond compliance
7. **AI Inventory:**
    - Purpose and use of AI system inventories
    - Risk scoring and resource allocation
8. **Stakeholder Engagement:**
    - Early engagement, diversity of perspectives
    - Leadership support strategies

Study Tips:

- Understand how to apply the NIST AI RMF functions in scenarios
- Practice calculating risk scores using severity × probability
- Know when each governance model (centralized/decentralized/hybrid) is appropriate

Narahari

Udemy Course Link : https://tinyurl.com/AIGPbyHari

- Understand the relationship between different risk management strategies
- Focus on practical application, not just definitions

Good luck with your AIGP certification!

Narahari