AIGP Module 5 Advanced Practice Test

Governing AI Deployment

30 Medium to Hard Questions (Exam Simulation)

Question 1

An organization is deciding where to deploy a healthcare AI system that processes sensitive patient data and requires significant computational power. According to Module 5, which deployment environment would provide the GREATEST control over data security but requires higher upfront hardware investment?

A. Cloud-based deployment
B. On-premise deployment
C. Edge deployment
D. Hybrid deployment

Correct Answer: B

Explanation:

- **Why B is correct:** The training states: "On-premise: hosting the model on your organization's servers and hardware - Greater control over deployment infrastructure (especially important if you handle sensitive data or are in a regulated sector); however, may require a greater upfront hardware investment." For sensitive healthcare data, on-premise provides maximum control.
- **Why A is incorrect:** Cloud-based has "latency and security risks due to a third party handling the data," making it less suitable when maximum control over sensitive data is required.
- **Why C is incorrect:** Edge deployment is "limited by edge device hardware, which can limit computational power." Healthcare AI requiring significant computation wouldn't suit edge deployment.
- **Why D is incorrect:** Hybrid isn't discussed in Module 5's deployment environment comparison.

Question 2

According to Module 5, what is Retrieval-Augmented Generation (RAG) in the context of generative AI deployment?

A. A security threat to AI systems
B. A process that optimizes LLM output by referencing a knowledge base beyond training data sources

Narahari

C. A type of edge deployment
D. A method for encrypting AI outputs

Correct Answer: B

Explanation:

- **Why B is correct:** The training defines RAG as: "a process that optimizes LLM output by referencing a knowledge base beyond training data sources." This allows generative AI to access current or domain-specific information not in training data.
- **Why A is incorrect:** RAG is a performance enhancement technique, not a security threat.
- **Why C is incorrect:** Edge deployment refers to hosting models on devices, not techniques for optimizing LLM outputs.
- **Why D is incorrect:** RAG is about enhancing outputs through external knowledge, not encryption.

Question 3

An organization discovers that its customer service AI, originally designed for handling inquiries, is being used by employees for making hiring decisions. According to Module 5, what governance concern does this scenario PRIMARILY illustrate?

A. Model drift from environmental changes
B. AI being used for a purpose it was not originally modeled for, requiring documentation of appropriate uses
C. Cloud deployment security issues
D. Red teaming exercises

Correct Answer: B

Explanation:

- **Why B is correct:** The training states: "Document appropriate uses of your AI to prevent use for a different purpose not intended in the creation of the AI. AI will not have the same appropriate factors when used for a new purpose" and identifies as a predictable risk: "AI being used for a purpose it was not originally modeled for." This is function creep or scope creep.
- **Why A is incorrect:** Model drift involves changes in data patterns over time affecting performance. This scenario describes misuse for unintended purposes, not drift from changing data conditions.
- **Why C is incorrect:** The issue isn't about deployment infrastructure security but about inappropriate use beyond the system's intended purpose and validation.
- **Why D is incorrect:** Red teaming is a testing methodology. The scenario describes inappropriate deployment scope, not a testing practice.

Narahari

Question 4

According to Module 5, what is model drift and what causes it?

A. Intentional updates to improve performance
B. When the relationship between input data and output predictions changes over time because conditions under which the model was trained no longer apply
C. A planned model replacement process
D. A type of security attack on AI systems

Correct Answer: B

Explanation:

- **Why B is correct:** The training states: "Model drift occurs when the relationship between input data and output predictions changes over time. This means the conditions under which the model was trained no longer apply, causing a decline in model performance." Example given: "a spam detection model that fails to recognize new types of spam when the nature of spam evolves."
- **Why A is incorrect:** Model drift is an unintended degradation, not intentional updates. Intentional improvements would be model retraining or updates, not drift.
- **Why C is incorrect:** Drift is unplanned degradation requiring intervention. Planned replacement is a governance activity in response to drift or other issues, not drift itself.
- **Why D is incorrect:** Drift is a natural phenomenon from changing conditions, not a security attack. Security attacks include model inversion, extraction, poisoning, and evasion.

Question 5

An organization wants to ensure its deployed AI system continues to perform effectively. According to Module 5, what is the recommended approach for detecting and addressing model drift?

A. Ignore performance changes and hope they resolve naturally
B. Define a baseline to measure future iterations, monitor performance, and retrain with fresh data reflecting current trends
C. Immediately shut down the system at the first sign of any change
D. Only check performance once per year

Correct Answer: B

Explanation:

- **Why B is correct:** The training states: "Continue to iterate the model to improve performance as the data changes. In addition, you can keep models up to date by retraining them with fresh data that reflects the most current trends and patterns. Define a baseline to measure future iterations of the model." This proactive approach maintains model relevance.

Narahari

- **Why A is incorrect:** The training emphasizes active monitoring and intervention: "Monitor and maintain the model to prevent model drift." Ignoring changes contradicts responsible governance.
- **Why C is incorrect:** While shutdown capability is necessary, immediately shutting down at any change is overreactive. The appropriate response is monitoring, assessment, and retraining as needed.
- **Why D is incorrect:** The training emphasizes "continuous monitoring" after deployment, not annual checks. Drift can occur rapidly, requiring more frequent monitoring.

Question 6

According to Module 5, what is red teaming in the context of AI deployment?

A. Marketing strategy for AI products
B. Structured testing efforts to find flaws and vulnerabilities in AI systems
C. A type of cloud deployment
D. A data encryption method

Correct Answer: B

Explanation:

- **Why B is correct:** The training defines red teaming as "structured testing efforts to find flaws and vulnerabilities" and states: "Conduct internal or external red teaming exercises for AI systems. These should also be done pre-deployment." Red teaming simulates adversarial attacks to identify weaknesses.
- **Why A is incorrect:** Red teaming is a security and robustness testing methodology, not a marketing strategy.
- **Why C is incorrect:** Cloud deployment is a hosting environment. Red teaming is a testing practice applicable to any deployment environment.
- **Why D is incorrect:** Red teaming is about adversarial testing, not encryption. It simulates attacks to find vulnerabilities, not protect data through encryption.

Question 7

According to Module 5, what are bug bounties and why are they recommended?

A. Financial rewards for purchasing AI systems
B. Programs to generate user engagement and extensive feedback by rewarding discovery of vulnerabilities
C. Discounts on cloud deployment services
D. Marketing campaigns for AI products

Correct Answer: B

Explanation:

Narahari

- **Why B is correct:** The training states: "Consider bug bashing/bug bounties to generate user engagement and extensive feedback." Bug bounties incentivize external security researchers and users to identify and report vulnerabilities, providing broader testing coverage than internal teams alone.
- **Why A is incorrect:** Bug bounties reward finding vulnerabilities in systems, not purchasing them. They're security tools, not sales incentives.
- **Why C is incorrect:** Bug bounties are security programs, not deployment service discounts. They're about identifying flaws, not reducing infrastructure costs.
- **Why D is incorrect:** Bug bounties serve security and quality assurance purposes, not marketing. While they may have PR benefits, their primary purpose is vulnerability identification.

Question 8

Under the EU AI Act, when must conformity assessments be performed for high-risk AI systems?

A. Only after deployment if problems occur
B. Before the system is put on the market, as well as over the life cycle of the system
C. Never—they are optional
D. Only for systems processing personal data

Correct Answer: B

Explanation:

- **Why B is correct:** The training states: "Are particular to high-risk AI systems and must take place before the system is put on the market, as well as over the life cycle of the system." This ensures compliance before release and continued compliance throughout operation.
- **Why A is incorrect:** Conformity assessments are proactive requirements before market entry, not reactive responses to post-deployment problems.
- **Why C is incorrect:** For high-risk AI systems under the EU AI Act, conformity assessments are mandatory, not optional. The training explicitly says they "Must be performed."
- **Why D is incorrect:** The training notes: "The requirement is not just for cases where personal information is being processed." High-risk determination is based on potential harm to health, safety, and fundamental rights, not just personal data processing.

Question 9

According to Module 5, what types of AI systems require conformity assessments under the EU AI Act?

Narahari

A. Only entertainment applications
B. High-risk systems including those used in recruitment, biometric identification, safety components, access to essential services, and critical infrastructure
C. Only systems developed by large companies
D. All AI systems regardless of risk level

Correct Answer: B

Explanation:

- **Why B is correct:** The training states conformity assessments "Apply to the use of AI in recruitment, biometric identification surveillance systems, safety components (e.g., medical devices), access to essential private and public services (e.g., creditworthiness, life insurance) and safety of critical infrastructure (e.g., energy, transport)." These are high-risk applications affecting fundamental rights.
- **Why A is incorrect:** Entertainment applications are typically low-risk and wouldn't require conformity assessments. The requirement applies to high-risk systems affecting rights, safety, and access to services.
- **Why C is incorrect:** The requirement is based on system risk level and application domain, not company size. Small companies deploying high-risk AI must also comply.
- **Why D is incorrect:** Conformity assessments are specifically for high-risk AI systems, not all systems. The EU AI Act uses a risk-based approach with different requirements for different risk levels.

Question 10

According to Module 5, what is the purpose of model cards in the deployment readiness process?

A. Marketing materials for selling AI systems
B. Standardized documentation of development, performance, intended use, and compliance to prepare for release
C. Financial budgeting tools
D. Only required for systems that fail testing

Correct Answer: B

Explanation:

- **Why B is correct:** The training states readiness includes "creating the model card" and notes: "Using model cards to document original purpose and new purposes" and mentions "Model Card Regulatory Check app automates regulatory compliance of AI systems based on accepted AI documentation tools like model cards." Model cards standardize essential information for deployment decisions and compliance.
- **Why A is incorrect:** Model cards are governance and compliance documentation, not marketing materials. They document technical details, limitations, and risks, which may not be promotional.

Narahari

- **Why C is incorrect:** Model cards document technical and governance information about the AI system itself, not financial budgets or cost information.
- **Why D is incorrect:** Model cards are standard documentation for all systems preparing for deployment, not just those that fail testing. They're proactive governance tools.

Question 11

According to Module 5, the FTC and EU AI Act require what regarding AI disclosure?

A. Disclosure is optional
B. Disclosure of AI use is required for almost any engagement
C. Only disclose if the AI makes errors
D. Only disclose to government regulators, not users

Correct Answer: B

Explanation:

- **Why B is correct:** The training states: "Under most laws, must disclose any AI being used. The FTC (U.S.) has been clear that disclosure is required for almost any engagement." This ensures transparency about when users interact with automated systems.
- **Why A is incorrect:** The training explicitly states disclosure is required under most laws, not optional. Both FTC and EU AI Act mandate disclosure.
- **Why C is incorrect:** Disclosure is required proactively for AI use, not just reactively when errors occur. Users need to know they're interacting with AI regardless of performance.
- **Why D is incorrect:** The training emphasizes disclosure to users: "Deployers have disclosure requirements in both directions: to users about the system and to the provider about incidents." Disclosure is required to both users and regulators.

Question 12

According to Module 5, in which sectors might there be additional disclosure requirements beyond just stating that AI is in use?

A. Only the technology sector
B. Financial services, health information, and education
C. Only manufacturing
D. No sectors have additional requirements

Correct Answer: B

Explanation:

- **Why B is correct:** The training states: "Areas like financial services, health information and education may have additional disclosure requirements, such as an obligation to disclose when a process is using AI. There can be requirements to disclose information

Narahari

beyond the fact that AI is in place (for example, if bias testing was done or the user has an opportunity to opt out of using AI)."

- **Why A is incorrect:** While technology companies deploy AI, the training specifically identifies financial services, healthcare, and education as sectors with enhanced disclosure requirements due to their sensitive nature and regulatory frameworks.
- **Why C is incorrect:** Manufacturing isn't specifically mentioned as having additional disclosure requirements. The training identifies financial services, healthcare, and education.
- **Why D is incorrect:** The training explicitly identifies multiple sectors with additional requirements beyond basic AI disclosure.

Question 13

According to Module 5, what are adverse action notices in the financial services industry?

A. Penalties for violating AI regulations
B. Notices provided when AI is used and a user is turned down for something like a loan
C. Marketing materials for financial products
D. Internal documents not shared with customers

Correct Answer: B

Explanation:

- **Why B is correct:** The training states: "The financial services industry uses adverse action notices when AI is used and a user is turned down for something, like a loan." These notices inform applicants of denials and provide transparency about automated decision-making.
- **Why A is incorrect:** Adverse action notices are customer communications about automated decisions affecting them, not penalties imposed on organizations for violations.
- **Why C is incorrect:** Adverse action notices inform customers of negative decisions, not promote products. They're regulatory compliance tools, not marketing.
- **Why D is incorrect:** Adverse action notices must be provided to affected customers, not kept internal. Transparency to users is the purpose.

Question 14

According to Module 5, what right do users have when an AI system produces an outcome not in their favor?

A. No rights—AI decisions are final
B. Appeal or redress options to exercise their legal protections
C. Only the right to delete their data
D. Automatic approval of their request

Narahari

Correct Answer: B

Explanation:

- **Why B is correct:** The training states: "There are other aspects of using an AI system that may require that AI be identified so users know how to exercise their rights or legal protections. For example, they may have appeal or redress options when an outcome is not in their favor and AI was used." This ensures human oversight and fairness.
- **Why A is incorrect:** Users have rights to challenge automated decisions. The training emphasizes appeal and redress mechanisms, not finality of AI decisions.
- **Why C is incorrect:** While data deletion rights may exist under privacy laws, the training specifically discusses appeal and redress rights related to adverse AI decisions, not just data deletion.
- **Why D is incorrect:** Appeal rights allow requests to be reviewed, not automatically approved. The right is to contest and have human review, not guaranteed reversal.

Question 15

According to Module 5, what capability should organizations have regarding AI system deactivation?

A. Only vendors can deactivate systems
B. The ability for a human to shut down an algorithm if it is not performing properly, especially for high-risk systems
C. Systems should never be deactivated once deployed
D. Only after regulatory approval

Correct Answer: B

Explanation:

- **Why B is correct:** The training states: "There should be the ability for a human to shut down an algorithm if it is not performing properly, especially if it has privacy impacts or the risks of using the algorithm are so great that there cannot be any room for error." This "kill switch" capability is essential for safety.
- **Why A is incorrect:** Organizations deploying AI must maintain control over shutdown capability, not depend solely on vendors. The training emphasizes "Ensure there is a procedure in place to deactivate a system."
- **Why C is incorrect:** The training explicitly requires deactivation procedures. Systems may need shutdown for performance issues, safety concerns, or legal requirements.
- **Why D is incorrect:** Organizations need immediate shutdown capability for emergencies, not delayed by regulatory approval processes. The kill switch must be available for urgent situations.

Question 16

Narahari

When evaluating vendor agreements for third-party AI, Module 5 recommends assessing data considerations. Which question is specifically mentioned?

A. Only the vendor's marketing budget
B. Whether they have legal rights to the data used, if personal data was minimized and de-identified, and how they will use data collected during operation
C. Only the vendor's office location
D. Only the vendor's stock price

Correct Answer: B

Explanation:

- **Why B is correct:** The training lists data considerations including: "Do they have legal rights to the data used? Was personal data minimized during collection and de-identified before being used for training or testing? If they will collect data your organization uses with the model, how will they use it?" These ensure lawful and privacy-protective data practices.
- **Why A is incorrect:** Marketing budgets aren't mentioned as relevant vendor agreement considerations. The focus is on technical, legal, and risk-related terms.
- **Why C is incorrect:** Vendor office location isn't mentioned as a key evaluation criterion. Data rights, security, and performance are the priorities.
- **Why D is incorrect:** Stock price isn't a vendor agreement evaluation criterion. The training focuses on legal compliance, security, bias mitigation, and performance.

Question 17

According to Module 5, what should organizations assess regarding bias when evaluating third-party AI vendors?

A. Bias is the vendor's responsibility only
B. What steps have been taken to minimize bias and whether they've established the AI doesn't lead to discrimination
C. Bias testing is unnecessary for vendor AI
D. Only check bias after deployment problems occur

Correct Answer: B

Explanation:

- **Why B is correct:** The training lists under vendor evaluation: "Potential for bias: What steps have been taken to minimize bias? Have they established that the AI function does not lead to statistical inaccuracies, bias or discrimination in results that apply to people?" Organizations must verify vendors' bias mitigation efforts.
- **Why A is incorrect:** While vendors are responsible for their systems, deploying organizations remain accountable and must assess vendor bias mitigation as part of due diligence.

Narahari

- **Why C is incorrect:** The training explicitly includes bias assessment as a key vendor evaluation criterion. Vendor AI requires the same scrutiny as internally developed systems.
- **Why D is incorrect:** Bias assessment should occur during vendor evaluation before deployment, not reactively after problems emerge. Proactive assessment is part of responsible procurement.

Question 18

According to Module 5, what should be clarified in vendor agreements regarding intellectual property and model outputs?

A. IP terms don't matter for AI
B. What is specified regarding intellectual property and ownership of model outputs, especially if the model generates content
C. Only the vendor owns all outputs
D. IP is never addressed in vendor contracts

Correct Answer: B

Explanation:

- **Why B is correct:** The training includes in vendor agreement evaluation: "If the model generates content, what is specified regarding intellectual property and model outputs?" This is crucial for generative AI where output ownership can be ambiguous and commercially significant.
- **Why A is incorrect:** IP terms are specifically mentioned as important vendor agreement considerations, especially for generative AI where output ownership affects commercial use rights.
- **Why C is incorrect:** Output ownership should be negotiated and clearly specified in contracts, not automatically assigned to vendors. Organizations may need rights to outputs for their business operations.
- **Why D is incorrect:** The training explicitly lists IP and output ownership as key terms to evaluate in vendor agreements, showing these must be addressed.

Question 19

According to Module 5, what unique challenges arise with proprietary AI models that organizations develop themselves?

A. No unique challenges—they're easier than third-party AI
B. Transparency issues from proprietary training data, output ownership questions, and potentially increased regulatory obligations
C. Proprietary models are always lower risk
D. Documentation isn't needed for proprietary models

Correct Answer: B

Narahari

Explanation:

- **Why B is correct:** The training states: "Transparency issues can arise from the nature of proprietary training data, or limitations in documentation. Ownership of output must be addressed if users generate derivative works. Consider regulatory/increased obligations for proprietary models in events like data breaches." Proprietary models may face enhanced scrutiny.
- **Why A is incorrect:** The training explicitly identifies "Unique Challenges with Proprietary AI," indicating they're not necessarily easier and have distinct governance considerations.
- **Why C is incorrect:** Proprietary models aren't automatically lower risk. The training notes they may face "increased regulatory obligations" and transparency challenges requiring careful management.
- **Why D is incorrect:** Documentation is essential for all AI systems. The training mentions "limitations in documentation" as a proprietary AI challenge to address, not eliminate documentation requirements.

Question 20

According to Module 5, how should organizations approach third-party AI risk management?

A. Trust vendors completely without oversight
B. Adapt procurement and vendor management processes to include AI-specific risk assessment, monitor integrated systems, and maintain clear policies
C. Avoid all third-party AI
D. Only review third-party AI once at procurement

Correct Answer: B

Explanation:

- **Why B is correct:** The training states: "Adapt procurement/vendor management processes to include AI-specific risk assessment and requirements. Clear AI-specific procurement and inventory policies are required. Monitor and assess performance and outputs of integrated third-party systems (e.g., using NIST AI Risk Management Framework)." This comprehensive approach manages third-party risks.
- **Why A is incorrect:** The training emphasizes organizations "may have limited visibility and control over third-party AI systems," requiring active oversight and monitoring, not blind trust.
- **Why C is incorrect:** Third-party AI can be used responsibly with proper governance. Avoiding all third-party AI is impractical and not recommended—the focus is on proper risk management.
- **Why D is incorrect:** The training emphasizes ongoing "Monitor and assess performance and outputs" of third-party systems, not just one-time procurement review. Continuous oversight is necessary.

Narahari

Question 21

According to Module 5, what questions should a readiness assessment answer before releasing an AI system into production?

A. Only marketing strategy questions
B. Is the model meeting performance, reliability, safety requirements? Are all requirements satisfied including documentation, explainability, and ethical standards?
C. Only financial return questions
D. Readiness assessments are unnecessary

Correct Answer: B

Explanation:

- **Why B is correct:** The training states: "Questions to answer: Is the model meeting performance, reliability, safety? Are all requirements satisfied (documentation, explainability, technical/ethical/measurable standards)?" A comprehensive readiness assessment ensures the system is prepared for deployment.
- **Why A is incorrect:** Readiness assessments focus on technical performance, safety, compliance, and ethical standards, not marketing strategy.
- **Why C is incorrect:** While financial considerations matter to business, readiness assessments evaluate technical performance, safety, compliance, and ethical requirements—not just financial returns.
- **Why D is incorrect:** The training explicitly discusses readiness assessments as important pre-deployment activities to ensure systems meet requirements before release.

Question 22

According to Module 5, what should continuous monitoring after deployment watch for?

A. Only positive user reviews
B. Deviations in accuracy, irregular decisions, and data drift
C. Only competitor AI systems
D. Monitoring isn't necessary after deployment

Correct Answer: B

Explanation:

- **Why B is correct:** The training states: "Continuous monitoring after deployment to watch for deviations in accuracy, irregular decisions, and data drift." These indicators reveal performance degradation or emerging issues requiring intervention.
- **Why A is incorrect:** Monitoring must detect problems, not just positive feedback. The training emphasizes watching for deviations, irregularities, and drift—negative indicators requiring response.

Narahari

- **Why C is incorrect:** Monitoring focuses on the deployed system's performance, not competitive intelligence about other companies' AI systems.
- **Why D is incorrect:** The training explicitly requires "Continuous monitoring after deployment," emphasizing it's not a one-time deployment activity but ongoing governance.

Question 23

According to Module 5, what should organizations maintain regarding model versions and changes?

A. Delete all previous versions immediately
B. Keep documentation including logs, versioning, and snapshots
C. Only keep the current version
D. Version control is unnecessary for AI

Correct Answer: B

Explanation:

- **Why B is correct:** The training states: "Keep documentation (model cards, logs, versioning/snapshots)." Maintaining version history supports troubleshooting, auditing, rollback capabilities, and accountability.
- **Why A is incorrect:** Deleting previous versions eliminates rollback options and audit trails. The training emphasizes keeping versioning and snapshots, not deleting them.
- **Why C is incorrect:** Version control requires maintaining multiple versions and their documentation. The training specifically mentions "versioning/snapshots" (plural), indicating multiple versions should be retained.
- **Why D is incorrect:** The training explicitly includes "versioning/snapshots" as essential documentation practices, showing version control is necessary for AI governance.

Question 24

According to Module 5, what is a challenger model and when should it be used?

A. A model that competes in marketing
B. A new model tested against the existing champion model to detect drift and performance changes
C. A model used only for training purposes
D. Challenger models are only used in development, never deployment

Correct Answer: B

Explanation:

- **Why B is correct:** The training states: "Create a 'challenger model' (a new model) to test and compare against the existing model ('champion model') to test for drift, unexpected results, etc." Challenger models provide independent validation of deployed systems.

Narahari

- **Why A is incorrect:** Challenger models are governance tools for performance validation, not marketing competition tools.
- **Why C is incorrect:** While models are trained, challenger models specifically serve to validate deployed champion models through comparative testing, not just for initial training.
- **Why D is incorrect:** The training discusses challenger models in the context of monitoring deployed systems: "Create a 'challenger model'..." appears in the post-deployment monitoring section, showing they're used after deployment.

Question 25

According to Module 5, when should organizations use current best practices like red teaming?

A. Only before deployment
B. Continuously when managing and monitoring internal and external risks, both pre-deployment and post-deployment
C. Never—red teaming is outdated
D. Only if problems occur

Correct Answer: B

Explanation:

- **Why B is correct:** The training states: "Best practices evolve constantly, so it is important to use current best practices when continuously managing and monitoring internal and external risks" and "Conduct internal or external red teaming exercises for AI systems. These should also be done pre-deployment." This indicates ongoing use both before and after deployment.
- **Why A is incorrect:** While red teaming should occur pre-deployment, the training emphasizes "continuously managing and monitoring," indicating ongoing use post-deployment as well.
- **Why C is incorrect:** The training explicitly recommends red teaming as a current best practice, not an outdated technique. It's a recognized method for finding vulnerabilities.
- **Why D is incorrect:** Red teaming is proactive security testing to find problems before they occur in production, not reactive testing after problems emerge.

Question 26

According to Module 5, what types of risks should organizations monitor as "predictable risks" after deployment?

A. Only completely unexpected risks
B. AI being used for a purpose it was not originally modeled for, and integration of new data
C. Only risks that have already occurred
D. Predictable risks don't need monitoring

Narahari

Correct Answer: B

Explanation:

- **Why B is correct:** The training states: "Some risks are more predictable. For example, AI being used for a purpose it was not originally modeled for" and mentions "new data integration" as foreseeable risks requiring monitoring.
- **Why A is incorrect:** The question asks about predictable risks. Completely unexpected risks are by definition unpredictable. The training distinguishes between predictable and unforeseen risks.
- **Why C is incorrect:** Predictable risks are those that can be foreseen and monitored proactively, not just historical risks that have already occurred.
- **Why D is incorrect:** The training explicitly discusses monitoring predictable risks, showing they require attention. Being predictable makes them easier to monitor systematically.

Question 27

According to Module 5, what should be included in an incident response plan for AI systems?

A. Only financial compensation procedures
B. Identifying and escalating issues, understanding reporting obligations, and enabling manual deactivation by a human
C. Only public relations strategies
D. Incident response plans are unnecessary for AI

Correct Answer: B

Explanation:

- **Why B is correct:** The training states: "Incident response includes identifying and escalating issues, understanding reporting obligations, and enabling manual deactivation ('kill switch') by a human if necessary." Comprehensive incident response covers detection, escalation, reporting, and emergency shutdown.
- **Why A is incorrect:** While compensation might be part of remediation, incident response focuses on identifying, escalating, reporting, and stopping problematic AI systems—not just financial settlements.
- **Why C is incorrect:** Public communications may be part of incident response, but the training emphasizes technical and operational elements: identification, escalation, reporting obligations, and deactivation capabilities.
- **Why D is incorrect:** The training explicitly discusses "Have an established incident response plan for when the system is not performing as intended," showing such plans are essential.

Narahari

Question 28

According to Module 5, what is the advantage of edge deployment for AI systems?

A. Unlimited computational power
B. May decrease latency and increase privacy
C. Always cheaper than other options
D. Easier to update than other deployment methods

Correct Answer: B

Explanation:

- **Why B is correct:** The training states: "Edge: hosting the model on edge devices like smartphones - May decrease latency and increase privacy." Processing data locally on devices reduces transmission delays and keeps sensitive data on the device.
- **Why A is incorrect:** The training notes edge deployment's disadvantage: "the model may be limited by edge device hardware, which can limit computational power." Edge devices have limited, not unlimited, power.
- **Why C is incorrect:** Cost isn't mentioned as an edge deployment advantage. The training focuses on latency and privacy benefits versus computational limitations.
- **Why D is incorrect:** The training doesn't discuss update ease. Edge deployment's advantages are latency reduction and privacy enhancement, not update simplicity.

Question 29

According to Module 5, what should organizations evaluate regarding model performance and monitoring in vendor agreements?

A. Performance standards don't need to be in contracts
B. What has been done to ensure model stability and prevent inaccurate outputs, and how the model will be monitored and maintained
C. Only the initial deployment, not ongoing performance
D. Vendors should not be required to monitor their models

Correct Answer: B

Explanation:

- **Why B is correct:** The training lists vendor agreement considerations including: "Model performance results: What has been done to ensure model stability and prevent inaccurate outputs? How will the model be monitored/maintained?" These contractual terms ensure ongoing accountability.
- **Why A is incorrect:** The training explicitly includes performance and monitoring in the list of key terms to evaluate in vendor agreements, showing they should be contractually addressed.

Narahari

- **Why C is incorrect:** The training asks "How will the model be monitored/maintained?" indicating ongoing performance and monitoring should be addressed, not just initial deployment.
- **Why D is incorrect:** Vendor monitoring and maintenance obligations are specifically mentioned as important contract terms to evaluate, showing vendors should be required to monitor their models.

Question 30

According to Module 5, which deployment environment advantage is "Easy to scale and reduced investment in hardware"?

A. On-premise deployment
B. Edge deployment
C. Cloud-based deployment
D. No deployment environment has this advantage

Correct Answer: C

Explanation:

- **Why C is correct:** The training states: "Cloud-based: a third-party cloud provider hosts the model and handles infrastructure - Easy to scale and reduced investment in hardware." Cloud providers offer elastic scaling without upfront hardware purchases.
- **Why A is incorrect:** On-premise deployment has the opposite characteristic: "may require a greater upfront hardware investment." Organizations must purchase and maintain their own infrastructure.
- **Why B is incorrect:** Edge deployment's advantages are "May decrease latency and increase privacy," not scalability or reduced hardware investment. Edge devices have limited resources.
- **Why D is incorrect:** Cloud-based deployment specifically offers these advantages as stated in the training.

End of Module 5 Advanced Practice Test

Scoring Guide:

- **27-30 correct:** Exceptional understanding - exam-ready
- **24-26 correct:** Strong understanding - minor review needed
- **21-23 correct:** Good foundation - review deployment environments and vendor management
- **18-20 correct:** Adequate knowledge - focus on monitoring and disclosure requirements
- **Below 18:** Review Module 5 thoroughly, especially deployment considerations and post-deployment governance

Narahari

Key Concepts to Master for Module 5:

1. **Deployment Environments:**
    - Cloud: Scalable, lower hardware cost, but latency/security concerns
    - On-premise: Greater control, suitable for sensitive data, higher cost
    - Edge: Lower latency, better privacy, limited computational power
2. **Deployment Techniques:**
    - Retrieval-Augmented Generation (RAG)
    - Fine-tuning vs. using models as-is
3. **Risk Management:**
    - Function creep (unintended purposes)
    - Model drift detection and mitigation
    - Challenger models
4. **Testing and Validation:**
    - Red teaming
    - Bug bounties
    - Conformity assessments (EU AI Act)
5. **Documentation:**
    - Model cards for deployment readiness
    - Version control and snapshots
    - Logs and audit trails
6. **Transparency Requirements:**
    - Disclosure obligations (FTC, EU AI Act)
    - Adverse action notices (financial services)
    - Sector-specific requirements (healthcare, education)
7. **User Rights:**
    - Right to appeal/redress
    - Human review of automated decisions
    - Opt-out options
8. **Operational Controls:**
    - Kill switches/deactivation procedures
    - Continuous monitoring

Narahari

    - Incident response plans

9. **Vendor Management:**
    - Data rights and minimization
    - Security and bias assessments
    - IP and output ownership
    - Performance monitoring obligations
10. **Ongoing Governance:**
    - Continuous monitoring for drift
    - Regular retraining
    - Performance baselines

Study Tips:

- Understand the tradeoffs between deployment environments
- Know what RAG is and when it's used
- Memorize disclosure requirements across jurisdictions
- Understand model drift causes and mitigation strategies
- Know the purpose and timing of conformity assessments
- Understand vendor agreement key terms
- Know when kill switches are necessary
- Practice distinguishing between predictable and unpredictable risks

Good luck with your AIGP certification!

Narahari